

Rishi More

USA | rmore2@jhu.edu | +1 (650)220-9450 | linkedin.com/in/rishimore102 | github.com/rishi-more-2003 | rishi-more-2003.github.io

Summary

ML engineer building **risk-controlled LLM inference, retrieval systems, and production-grade ML pipelines**. Co-first author of an **ICML 2026** paper on distribution-free risk control for reducing test-time compute under calibrated error guarantees. Experienced in turning research prototypes into low-latency ML pipelines, evaluation harnesses, and production-facing systems.

Experience

JHU Data Science and AI Institute (DSAI)

May 2025 – Present

Graduate Research Assistant (Advisors: Prof. Daniel Khashabi, Prof. Eric Nalisnick)

Baltimore, MD

- **Co-first author** of *Conformal Thinking (ICML 2026)*: distribution-free risk control that sets dual confidence thresholds for adaptive reasoning, cutting LLM inference cost while bounding error at a user-specified risk level.
- Built a vLLM-based calibration/evaluation pipeline across **4** reasoning models and **4** benchmarks (**2.7K+** problems), maintained target error rates under dataset/length shift using finite-sample (UCB) calibration.
- Designed dual-threshold stopping rules, including a parametric lower threshold and ensembled halt policy, reducing reasoning tokens by **12%** at matched accuracy with reusable evaluation tooling.

JHU Center for Language and Speech Processing (CLSP)

January 2025 – December 2025

Graduate Researcher (Advisor: Prof. Jason Eisner)

Baltimore, MD

- Built a temporal-aware RAG pipeline over **10K+** non-stationary documents using time-decay weighting, metadata-aware indexing, and temporally scoped retrieval.
- Improved retrieval precision by **23%** over dense baselines in high-volatility knowledge domains by prioritizing time-relevant evidence.
- Reduced false retrievals by **31%** using reference-free consistency scoring to detect stale or hallucination-prone retrieved context.

Mastek

November 2022 – July 2023

Machine Learning Engineer Intern

Mumbai, India

- Built and deployed an on-prem **edge** voice authentication system at **<100ms p99** latency and **96.6%** accuracy, using a lightweight Siamese mel-spectrogram model for retraining-free enrollment under edge-compute limits.
- Engineered a TensorFlow pipeline over **1M+** voice samples with spectral-gating denoising for robust feature extraction.
- Improved authentication reliability **20%** in high-noise production environments via preprocessing and feature-extraction upgrades.

Projects

Post-Training Evaluation & Reward-Hacking Detector | *Python, PyTorch, DPO, LLM-as-Judge, Calibration* | GitHub

- Built an evaluation harness for **6** post-training failure modes—reward hacking, sycophancy, verbosity bias, false confidence, safety drift, and distribution shift across **4** Qwen3-8B variants: Base, SFT, DPO, and Instruct.
- Evaluated **142** adversarial prompts with debiased pairwise LLM judging, position-swapped comparisons, and log-prob calibration, flagging **10%** of DPO preference wins that degraded factuality.
- Diagnosed optimization-stage regressions using ECE, Brier score, AUROC, and bootstrap confidence intervals, finding SFT raised harmful compliance to **12.5%** and DPO increased sycophancy from **34%** to **45%**.

Asymmetric Relation Dynamics in Language Agents | *Python, LLM Agents, LoRA, asyncio, BM25* | GitHub

- Built an async multi-agent training system (Tinker API): a frozen **Qwen3-235B** caregiver teaching LoRA-adapted **Qwen3-8B** agents via a four-module memory and salience-gated BM25 retrieval.
- Reached **100%** training success with **~25%** fewer turns (6.1 vs 8.4) over Solo/Peer baselines across **4** conditions $\times$ **3** seeds.
- Isolated scaffolding dependency as the transfer bottleneck via ablations; fully reproducible via a single-command CLI.

Label-Preserving Domain Adaptation via KL-Regularized Optimal Transport | *Python, Sinkhorn, K-Means, t-SNE* | GitHub

- Extended Sinkhorn optimal transport with KL information loss and intra-cluster consistency regularization for unsupervised label-preserving domain adaptation.
- Reduced transport computation from sample-level $O(n_s^2)$ to centroid-level $O(n_c^2)$ by clustering source and target distributions before alignment.
- Improved 1-NN target accuracy over Sinkhorn by **+12.6pp** on Two-Moons (99.8%) and **+5.3pp** on MNIST→USPS (82.7%).

Education

Johns Hopkins University

May 2026

M.S.E. in Computer Science

GPA: 4.0/4.0

- **Coursework:** Theory of Replicable ML, Human-in-the-Loop ML, Non-Stationary Environments.

University of Mumbai

July 2024

B.Tech. in Computer Science, Honors in AI/ML

CGPA: 9.85/10

Technical Skills

Languages: Python, C++, Java, C

ML/LLM Systems: PyTorch, TensorFlow, Hugging Face, LoRA, DPO, RAG, calibration, conformal/risk control, Scikit-learn

Inference & MLOps: vLLM, Docker, Kubernetes, AWS, Apache Airflow, Apache Spark

Data: SQL, PostgreSQL, MongoDB, Pandas, NumPy