

Rishi More

Boston, MA | rishi.more102@gmail.com | +1 (650) 220-9450 | linkedin.com/in/rishimore102 | github.com/rishi-more-2003 | rishi-more-2003.github.io

Summary

ML systems engineer focused on GPU-accelerated LLM inference, quantization, and reliable adaptive serving. **ICML 2026 co-first author** with experience building CUDA/Triton kernels and vLLM evaluation infrastructure.

Experience

InterSystems

Core Development Program, AI Experience

July 2026 – September 2026

Boston, Massachusetts

- Developed **LLM-backed clinical features** for **TrakCare**, a unified electronic health record deployed across **29 countries**.
- Integrated an **AI-generated patient summary column** into clinician workflows, consolidating patient histories into a single overview to simplify chart review and reduce manual navigation across clinical records.

Johns Hopkins University Data Science and AI Institute (DSAI)

Graduate Research Assistant, LLM Inference Optimization

May 2025 – Present

Baltimore, Maryland

- Co-first authored *Conformal Thinking* (**ICML 2026 main track**); designed **dual-threshold stopping** to exit on confident answers or stalled reasoning, reducing reasoning tokens by **12% at matched accuracy**.
- Built a calibration and evaluation pipeline using vLLM across **4 reasoning models, 4 benchmarks, and 2.7K+ problems**; captured intermediate answers, confidence trajectories, and token counts to compare stopping policies.
- Implemented **UCB-calibrated signal-threshold selection**, choosing the lowest-token policy satisfying a user-set risk budget; benchmarked against naive threshold tuning and stress-tested false-positive and premature-stop risk under dataset and response-length shifts.

Mastek

Machine Learning Engineer Intern

November 2022 – July 2023

Mumbai, India

- Deployed an on-premises edge voice-authentication system with **<100 ms p99 latency** and **96.6% accuracy**, using a lightweight Siamese mel-spectrogram model for retraining-free enrollment.
- Built a TensorFlow pipeline over **1M+ voice samples**; spectral-gating denoising, augmentation, and robust feature extraction improved authentication reliability by **20%** in high-noise conditions.

Projects

Conformal Serve: Risk-Controlled Adaptive LLM Serving | vLLM, FastAPI, Docker Compose, Prometheus

- Built an engine-agnostic adaptive-stopping core shared by deterministic mock-engine CI and vLLM's AsyncLLMEngine; achieved **100% policy branch coverage** with `mypy --strict` and exposed client-set risk α through an OpenAI-compatible API.
- On Qwen2.5-1.5B-Instruct and an 8GB RTX 4070, reduced reasoning tokens **64.7% at $\alpha = 0.10$** and **71.4% at $\alpha = 0.30$** while maintaining realized error $\leq \alpha$ on small held-out arithmetic sets.
- Calibrated per- α thresholds offline by distribution-free risk control into immutable, content-addressed artifacts; shipped Docker Compose with Prometheus/Grafana and a one-command benchmark reporting **95% Clopper–Pearson error bounds**.

INT4 KV-Cache Quantization & Fused Triton Attention for Long-Context Decode | Triton, CUDA, PyTorch

- Built a correctness-gated fused Triton flash-decode dequant-attention kernel, cutting 32K-context attention latency **13.7 $\times$ (7.54 to 0.55 ms/layer)** versus PyTorch FP16 GQA SDPA while maintaining **0.999 cosine similarity** to the dequantized oracle.
- Reduced KV storage from **28.7 KB to 8.3 KB per token (3.47 $\times$)**, raising estimated 8GB KV capacity from **$\sim 154K$ to $\sim 534K$ tokens**; retained **100% passkey retrieval** through 32K with only **+1.0% WikiText perplexity**.
- Diagnosed per-token-K failure as outlier-channel collapse (attention cosine **0.08 vs. 0.98**) and built an RCPS controller certifying **$\sim 50\%$ cache eviction at 19% realized loss** under a **30% risk budget**.

Decode Roofline: Kernel Profiling & Fused INT4 Dequant-GEMV | CUDA C++, Nsight

- Profiled Qwen2.5-1.5B batch-1 decode with Nsight; attributed **$\sim 81\%$ of GPU-kernel time** to weight GEMVs and measured dominant MLP projections at **~ 247 GB/s** against a **~ 250 GB/s** DRAM ceiling.
- Built a correctness-gated fused CUDA INT4 group-dequantization + GEMV kernel reaching **~ 223 GB/s**; accelerated batch-1 MLP projections by **2.38 $\times$** versus FP16 GEMM and mapped the crossover to batch 2 as weight reuse made GEMM faster.

Education

Johns Hopkins University

M.S.E. in Computer Science

May 2026

GPA: 3.9/4.0

University of Mumbai

B.Tech. in Computer Science, Honors in AI/ML

July 2024

CGPA: 9.85/10

Technical Skills

Languages: Python, C++, CUDA C++, C, SQL, Bash

ML/LLM Systems: PyTorch, Triton, vLLM, Hugging Face Transformers, TensorFlow, LoRA, DPO, RAG

Performance & Infrastructure: Nsight Compute/Systems, Linux, Git, Docker, GitHub Actions, AWS

Methods: Quantization, kernel fusion, roofline analysis, LLM evaluation, conformal risk control, calibration